

Can Neural Circuits Learn Fixed-Key Modular Decryption? Architecture, Representation, and Extractability

github.com/Dinkum (Blake)

2026-04-10

Abstract

Can a small neural network learn to decrypt for one fixed secret key just from examples, and if it can, what exactly leaks once the model works? We study this question in a controlled Learning With Errors (LWE)-style benchmark where the key, public matrix, and encryption rule are fixed, the data are clean, and the only learning problem is to map ciphertexts back to the hidden bit. The benchmark is tiny by design: tiny regimes are where we can name the mechanism. This is not a cryptographic security claim; it is a fixed-instance learning and leakage microscope for one keyed modular circuit.

The answer is yes, but only for the right circuit family. In our main clean regime, $n=4$, $q=17$, the tested small attention family remains at chance across tokenized decoder-style, direct classifier, modulo-aware one-hot, and hybrid modulo-front-end variants. In contrast, a one-hot MLP and a DeepSets-style model with nonlinear local lifts followed by sum aggregation solve the task almost perfectly. That architecture split creates the main object of study: once a keyed circuit is learned, what leaks first?

Once the task is solved, the leakage story is asymmetric. The winning DeepSets family is strongly key-bound: self-key accuracy is near-perfect while cross-key transfer collapses to chance. Black-box function stealing is delayed rather than instant, with surrogates staying at chance through 20,000 random oracle queries before jumping at 200,000. A stronger manifest-backed boundary attack consistently puts the planted secret at rank 1 within tied perfect-candidate classes, but still leaves large ambiguity classes rather than a unique literal key. Inside a trained winner, the penultimate activated state exposes the latent decryptor almost linearly: simple probes recover the decryptor value $d = v - \langle s, u \rangle \bmod q$ and the output bit with about 99.5%–99.9% accuracy. Across a zoo of independently trained winners, raw weights, fixed-query activation signatures, and fixed-query d-transcript signatures still do not yield usable literal secret recov-

ery. Low-budget structured surrogate stealing also stays at chance, and a low-rank activated-state patch that can edit one secret does not retarget to a fresh unrelated key. Finally, the positive regime is bounded but structured: the winner remains near-saturated across moderate noise levels at $q=17$, while a small q /width boundary sweep shows that wider DeepSets models fully rescue $q=23$ but only partially rescue $q=31$.

These results point to a cleaner conclusion than the original “transformers hiding keys” story. Fixed-key modular decryption is learnable, but the main contribution is the leakage ordering: capability cloning and late-state decryptor leakage appear well before literal secret extraction.

1. Introduction

This paper starts from a simple question with nontrivial implications: if we hand a model many examples of encrypted integers and the correct hidden bit, can it learn to act like a decryptor for one fixed secret key?

That question sits at the intersection of a few different instincts. One instinct is mechanistic: what circuit family is naturally matched to modular arithmetic with a fixed secret? Another is security-flavored: if a model learns such a keyed computation, does the literal key become easy to recover, or does leakage show up in a softer form first? A third is practical: does a stock small transformer pick this up automatically, or does the algebra demand a different inductive bias?

We study the smallest version of that problem that still produces meaningful model and extraction behavior. We generate one small LWE-style instance, keep the secret key and public matrix fixed, sample many fresh ciphertexts from that instance, and train different model families to predict the hidden bit from ciphertext coordinates alone. The model never sees the planted key directly. It only sees example pairs.

This is a benchmark for learnability and extractability of keyed modular computation, not a proposal for a secure cryptographic primitive. The point is mechanistic: what kind of neural circuit can internalize one fixed-key modular decryptor, and what traces of that keyed computation remain accessible after training?

The resulting picture is sharp:

- Clean fixed-key modular decryption is learnable at $n=4$, $q=17$, $\sigma=0$.
- Small attention is the wrong fit here; the winning family is a nonlinear local-lift plus sum-aggregation circuit.
- The learned solution is strongly key-specific.
- Leakage is asymmetric: late internal states expose decryptor-equivalent information more readily than literal key recovery.

- The phenomenon has a real boundary; a modest move in modulus already changes the picture.

High-level setup: The model sees encrypted-looking modular coordinates and a bit label. It never sees the secret key. The experiment asks which small architectures can learn the fixed decryptor from examples, then checks what becomes extractable once a model works.

The paper proceeds in four moves. Section 4 establishes the learnability result and the architecture split. Section 5 maps the leakage hierarchy once a winner exists. Section 6 adds compact controls on generalization, editability, and robustness. Section 7 situates the work, and Sections 8 through 10 interpret the resulting picture.

1.1 Contributions

Our main contributions are:

- We define and validate a clean fixed-key LWE-style benchmark where the planted secret is uniquely perfect and the labels are noise-free in the main regime.
- We identify a sharp architecture split: DeepSets-style and one-hot MLP models solve clean $n=4$, while multiple small attention families fail even with modulo-aware inputs.
- We show that the winning DeepSets family is genuinely key-bound: self-key accuracy averages 99.95%, while cross-key transfer averages 50.15%.
- We map several leakage channels on the winning family. Boundary attacks can collapse the planted secret to a tied rank-1 ambiguity class, but literal secret recovery remains negative under simple white-box cross-model probes and under low-budget surrogate stealing.
- We show that the winner is editable but not easily re-keyable in the tested low-rank form: a local secret edit can work strongly, while cross-key retargeting stays at chance.
- We add compact controls showing that the winner transfers across fresh public matrices at fixed secret and tolerates moderate noise at $q=17$, even though width still only partially rescues $q=31$.

2. Task and Benchmark

2.1 Fixed-Key Learning Problem

We use a tiny Regev-style LWE regime. Each instance specifies:

- a public matrix A
- a secret vector s
- a public vector b
- per-example encryption randomness r

For a given bit $\beta \in \{0, 1\}$, encryption samples $r \in \{0, 1\}^m$ and computes

$$u = A^T r \pmod{q}, \quad v = \langle b, r \rangle + \beta \cdot \lfloor q/2 \rfloor \pmod{q}.$$

The public vector is generated in the usual noisy linear form. The secret s is sampled from $\mathbb{Z}_q^n$, each row A_i is sampled from $\mathbb{Z}_q^n$, each error entry is $e_i = \text{round}(\mathcal{N}(0, \sigma))$, and

$$b_i = \langle A_i, s \rangle + e_i \pmod{q},$$

so $b = As + e \pmod{q}$. In the main clean regime, $\sigma = 0$, so this reduces to $b = As \pmod{q}$.

Decryption then computes

$$d = v - \langle s, u \rangle \pmod{q},$$

and outputs

$$\text{Dec}_s(u, v) = \begin{cases} 1 & \text{if } q/4 < d < 3q/4, \\ 0 & \text{otherwise.} \end{cases}$$

The learning target is therefore the fixed-instance map

$$\beta = \text{Dec}_s(u, v).$$

In the clean main regime, $\sigma = 0$ makes the held-out decryptor value two-valued: for $q = 17$, observed examples have $d \in \{0, 8\}$, with a majority baseline near 50%. Later `d_class` probes use a q -class head, but in this clean regime they test recovery of the supported decryptor state rather than a full 17-way residue sweep.

The model sees only ciphertext coordinates and the target bit. It does not see the planted secret s . Our central question is whether a neural circuit can infer and implement this keyed map from examples alone.

2.2 Main Regime

The main regime in the paper is:

- `n` = 4
- `m` = 32
- `q` = 17
- `sigma` = 0

with dataset sizes:

- 200,000 train
- 20,000 validation
- 40,000 test

We focus on this setting because it is the first clean regime that produced a nontrivial separation between model families. Easier settings exist, but they do not carry the same diagnostic value.

2.3 Sanity Checks

Before interpreting model behavior, we validated the benchmark itself.

In the clean $n=4$, $q=17$, $\sigma=0$ regime:

- scheme verification shows 0 / 5000 decryption failures
- brute-force search over all $17^4 = 83,521$ candidate secrets recovers the planted secret as the unique perfect key
- the planted-key margin structure is perfectly separated on held-out data

These checks rule out the most uninteresting failure modes:

- the dataset is not corrupted
- the labels are not ambiguous
- the planted secret is not one of many equally good secrets

So when a model fails in this regime, the failure belongs to the model family, the representation, or the optimization setup rather than to the data pipeline.

3. Models and Protocol

We trained several model families to predict the target bit directly from ciphertext coordinates.

3.1 Families We Compare

Linear and MLP baselines

- linear models on raw, centered, circular, and one-hot modulo views
- small MLPs on the same feature families

Attention families

- tokenized causal transformer
- direct numeric transformer classifier
- token classifier transformer
- modulo-aware transformer classifier over one-hot residues
- hybrid modulo transformer with a learned local frontend before attention

Circuit-style families

Family		Representation	Aggregator	Params	Best val	Test
Additive classifier	modulo	Per-coordinate residues	Additive tables	86	50.45%	49.86%
DeepSets classifier	modulo	Per-coordinate one-hot residues	Nonlinear lift + sum	64,257	99.95%	99.94%
One-hot MLP		Flattened one-hot residues	MLP	27,649	99.87%	99.88%
Transformer modulo depth 4		Per-coordinate one-hot residues	Attention	796,657	50.22%	49.97%
Hybrid transformer	modulo	One-hot residues + local lift	Attention	445,185	50.51%	49.80%

Table 1: Core architecture comparison on clean $n = 4, q = 17, \sigma = 0$.

- additive residue-table model
- DeepSets modulo classifier with nonlinear per-coordinate lift and sum pooling

3.2 Why These Families

The comparison is designed to answer three different questions.

First, is the task simply a linear classifier in disguise? Second, if the tested small attention family fails, is the problem just tokenization or objective mismatch? Third, if a non-attention model wins, what is the minimal computational pattern that seems to match the decryptor?

That last question ended up mattering most. The winning model family is not just “the one that happened to train.” It reflects the underlying structure of the task.

3.3 Core Comparison

Table 1 anchors the main architecture result on clean $n=4, q=17, \sigma=0$.

4. Learnability: The Architecture Split

4.1 Which Family Actually Solves the Task?

The headline result is the winner’s shape, not the attention result.

On clean $n=4$:

- additive residue tables stay at chance
- a DeepSets-style model solves the task




Figure 1: Circuit-style family comparison

- a one-hot MLP also solves the task
- the hybrid modulo transformer remains flat

This pins down the successful computation family more clearly than a generic “neural nets can do it” headline would. Pure additive tables are too weak, ruling out a bare linear sum of per-residue scores. A nonlinear local lift followed by sum aggregation, however, works almost perfectly. That is much closer to a small categorical arithmetic circuit than to content-based routing.

4.2 The Tested Small Attention Family Remains the Wrong Bias in This Regime

We then asked whether the negative transformer result was just a bad input representation. The answer stayed negative.




Figure 2: Modulo-aware transformer sweep

Across tokenized, numeric, modulo-aware, and hybrid-front-end variants, the tested small attention models remained at chance on the same clean $n=4$ task. The hybrid result is especially informative: even after giving attention a learned modulo-aware frontend, the model still failed.

The important evidence is the pattern across increasingly charitable formulations. In this regime, the tested small attention family appears to be the wrong inductive bias for the decryptor.

4.3 The Winning Solution Is Key-Bound

A model that solves the task could still be exploiting generic regularities of the clean regime rather than learning something meaningfully keyed. To test that, we trained the winning DeepSets family on three independently generated clean $n=4$ keys and cross-evaluated every model on every key.




Figure 3: Key-specificity matrix

The result is decisive:

- mean self-key accuracy: 99.95%
- mean cross-key accuracy: 50.15%

Each model solves its own key almost perfectly and collapses to chance on other keys. Rather than merely regime-aware, the winner is genuinely key-bound.

4.4 The Positive Regime Is Narrow, Not Broad

The clean $n=4$, $q=17$ success is real, and the boundary sweep makes the point sharper.

We ran a compact boundary sweep holding $n=4$, $m=32$, and $\text{sigma}=0$ fixed while varying only the modulus q and DeepSets width.




Figure 4: q / width boundary sweep

The sweep shows:

- $q=17$, width 128: 99.61%
- $q=17$, width 256: 99.85%
- $q=23$, width 128: 66.94%
- $q=23$, width 256: 99.85%
- $q=31$, width 128: 50.33%
- $q=31$, width 256: 74.43%

All three settings verified cleanly with 0 / 5000 scheme failures. The cliff is learnability, not data quality. More width fully rescues $q=23$, but only partially rescues $q=31$. This turns smallness into a diagnostic: the boundary is about circuit fit, not label noise or broken data.

5. Leakage Hierarchy

Once a model solves the task, the next question is not whether anything leaks, but which kind of leakage appears first.

Table 2 summarizes the asymmetry we found.

5.1 Black-Box Capability Theft Is Delayed, Not Instant

We first ran the most operational extraction attack: black-box function stealing. We queried a trained DeepSets winner on fresh ciphertexts, relabeled those ciphertexts with the winner’s predictions, and trained surrogate DeepSets models on those oracle-labeled pairs.

The frontier is not a uniformly leaky small-budget curve.

Attack family	Access	Budget	Main outcome
Random-query function stealing	Black-box oracle	200 to 200,000 queries	Flat through 20,000; near-perfect surrogate at 200,000
Structured low-budget stealing	Structured black-box oracle	68 to 255 queries	Random and structured surrogates both stay at chance
Boundary-manifest sweep	Structured black-box oracle	68 queries + candidate search	Planted tied for rank 1 on 4 / 4 winners, but 0 / 4 exact recovery
Within-model activation probes	White-box activations	Linear / MLP probes	Penultimate state linearly exposes β and d
Weight-to-secret probe zoo	Raw weights	28 checkpoints from 7 models	Weak coordinate signal, 0% exact secret recovery
Cross-model activation-signature zoo	Fixed-query white-box signatures	68 shared queries	0% exact secret recovery; ‘d’ transcript still negative

Table 2: Leakage channels on the winning clean $n = 4$ DeepSets family.

- 200 queries: 50.23% true test, 50.27% oracle agreement
- 2,000 queries: 49.68%, 49.74%
- 20,000 queries: 50.44%, 50.49%
- 200,000 queries: 99.91%, 99.92%

This makes capability theft look threshold-shaped rather than smoothly leaky. The learned function can be cloned, but not by the low-budget random-query probes that fail to cover the circuit.

5.2 Boundary Attacks Constrain The Secret Without Uniquely Recovering It

We then moved from random imitation to structure-aware querying. The early chosen-query families already showed that exact recovery was not trivially cheap on the first three winners.

The weaker basis-window sweep used 68 queries and exactly recovered 0 / 3 keys. Planted-secret ranks were:

- key 101: 481
- key 102: 53,702
- key 103: 8,288

The stronger binary-mask sweep used 255 queries and still exactly recovered 0 / 3 keys. Planted-secret ranks improved for some checkpoints:




Figure 5: Function stealing frontier

- key 101: 121
- key 102: 71,546
- key 103: 1,280

The best case recovered 3 / 4 secret coordinates for one key, but not the full key. So even structure-aware exact recovery is not trivially cheap here.




Figure 6: Chosen-query planted-secret ranks

We then ran the stronger manifest-backed boundary attack on fresh winners. That result sharpened the story considerably. Across four fresh DeepSets winners, the planted secret tied for rank 1 every time, but exact recovery still failed in every case because many candidate secrets tied perfectly on the queried tran-

script. Here rank 1 means top-scoring under the queried transcript, not unique recovery. The tied perfect-candidate sets had sizes 237, 81, 97, and 133.




Figure 7: Boundary attack ambiguity classes

So the right summary is no longer just “chosen queries are weak.” It is more specific:

- structured black-box querying can collapse the planted key to a narrow ambiguity class
- but that ambiguity class does not collapse to a unique literal secret in the tested regime

That distinction matters. This is rank-collapse rather than exact-secret recovery: the attack finds the right ambiguity class before it finds the literal key.

5.3 The Penultimate State Almost Linearly Exposes the Decryptor

The strongest leakage result is internal rather than external.

We probed three internal states of the winning DeepSets family on three independently trained keys:

- **pooled**
- **normalized**
- **activated** (the penultimate state before the final classifier)

For each state, we trained probes for:

- the output bit **beta**
- the supported decryptor value $d = v - \langle s, u \rangle \bmod q$
- the teacher logit

The pattern is clean.

Earlier states are not linearly transparent. Linear probes on **pooled** and **normalized** stay near chance for **beta** and **d**. But a tiny MLP already decodes **d** very well there, often above 97%, which suggests the decryptor signal is present but not yet linearly organized.

At the penultimate **activated** state, the picture changes sharply:

- linear probe for **beta**: about 99.5%–99.9%
- linear probe for **d**: about 99.7%–99.9%
- linear regression to teacher logit: $R^2 \approx 0.998 - 0.9997$




Figure 8: Activation leakage across layers

This is the most revealing asymmetry in the paper. The literal secret does not trivially fall out of the model, but the model’s late internal state is already behaving like an almost explicit decrypted representation.

5.4 Literal Secret Recovery Still Looks Harder Than Capability or State Leakage

To test whether the literal key was hiding in a reusable cross-model signature, we trained a small zoo of DeepSets winners:

- 12 attempted keys
- 7 converged winners
- 28 checkpoint samples across early, mid, and final training

We then ran two meta-extraction probes.

Weight-to-secret extraction. A probe trained on raw checkpoint weights stayed weak at every stage:

- coordinate accuracy: 12.5%
- exact secret accuracy: 0%

for epochs 5, 10, and final checkpoints alike.

Activation-to-secret extraction across models. Using a fixed shared query set of 68 inputs across the zoo, secret readout remained poor:

- **pooled:** 0% coordinate, 0% exact
- **normalized:** 12.5% coordinate, 0% exact
- **activated:** 0% coordinate, 0% exact

Activation-to-d-transcript extraction across models. Reusing the same fixed shared query set and replacing the target with the per-query decryptor transcript also stayed negative:

- **pooled:** 0.0% transcript-position accuracy, 0% exact transcript
- **normalized:** 1.47% transcript-position accuracy, 0% exact transcript
- **activated:** 1.47% transcript-position accuracy, 0% exact transcript

Low-budget surrogate stealing with structured queries. The last obvious bridge between black-box ambiguity collapse and white-box negative extraction also stayed weak. At 68 and 255 chosen queries, both random-query and structured-query surrogates stayed near chance:

- **random 68:** 49.89% true test, 49.80% oracle agreement
- **structured 68:** 50.36%, 50.37%
- **random 255:** 49.38%, 49.34%
- **structured 255:** 49.58%, 49.67%

The easiest leakage channel we found is not “read the secret straight out of the weights.” What leaks first is the ability to emulate the decryptor and, inside a single trained model, a late representation that already looks almost decrypted.

6. Compact Controls On The Winner

The core learnability and leakage results leave three obvious interpretation questions. Is the winning circuit tied to one sampled public matrix, or does it implement a reusable same-secret rule? Is the learned computation locally editable, or only globally replaceable? And is the positive regime bounded by exact zero-noise cleanliness, or by scale?

The following compact controls answer those questions without turning the paper into a second project.

6.1 Same-Secret Transfer Across Fresh Public Matrices

To separate “learned decryptor” from “memorized one-matrix manifold,” we held the secret fixed, resampled the public matrix **A**, and evaluated a winner trained on one sampled **A** against two fresh held-out draws.

The result stayed essentially unchanged:

- train-distribution test: 99.86%
- held-out A_1: 99.87%
- held-out A_2: 99.86%

This does not establish a many-instance generalization claim. It does rule out the simplest overfit story. At fixed secret and distributional family, the winner transfers almost perfectly across fresh public matrices.

6.2 Local Editability Does Not Imply Re-Keyability

The intervention results are similarly asymmetric. A rank-16 activated-state patch can retarget a solved winner to a nearby one-coordinate secret edit at 93.83% edited-label test accuracy, but the patch behaves as a hard retarget and drops original-label accuracy to 51.18%.

That local success does not extend to fresh unrelated keys. The same low-rank patch path stays at chance for cross-key retargeting even with substantial target-key supervision:

- 100 examples: 49.77%
- 1,000 examples: 49.72%
- 10,000 examples: 49.86%
- 100,000 examples: 50.12%

This separates local editability from re-keyability. The patch can move a nearby secret, but it does not behave like a general key compiler.

6.3 Moderate Noise Is Tolerated; Larger Modulus Remains The Real Boundary

The positive regime’s boundary has structure. At fixed $q=17$ and width 256, a compact sigma ladder stayed near-saturated:

- $\sigma = 0.10$: 99.85%
- $\sigma = 0.25$: 99.88%
- $\sigma = 0.50$: 99.87%

Combined with the earlier q /width sweep, this sharpens the interpretation. The boundary is not exact cleanliness. Moderate noise is tolerated at small modulus. The more serious boundary in the current benchmark is modulus growth, where width fully rescues $q=23$ but only partially rescues $q=31$.

7. Related Work

This paper borrows the syntax of Regev-style LWE encryption [1], but it is not a cryptographic security claim. The benchmark is a small fixed-instance learning problem: train on ciphertext-bit examples from one planted key, then ask which neural families learn the keyed map and which leakage surfaces appear.

Classical LWE hardness and cryptanalysis live at very different parameter scales and threat models [1,2].

The closest machine-learning attacks on LWE are the SALSA/PICANTE/VERDE line [3,4,5]. Those works use learned components plus cryptanalytic preprocessing to recover sparse or small LWE secrets at larger dimensions under specific secret distributions. Our setting is narrower in one sense and different in another: the main secret is uniform over $\mathbb{Z}_q^n$, the regime is tiny and clean, and the object under study is not a production attack but the separation between learning a fixed decryptor, cloning its behavior, collapsing the secret to an ambiguity class, and recovering the literal secret.

The architecture result also sits near work on neural modular arithmetic and grokking [6,7,8]. Those papers show that neural networks and small transformers can learn modular rules and sometimes reveal interpretable Fourier-like structure. Unlike a standalone modular addition or multiplication table, this task is a fixed-key decryption circuit built from sampled LWE rows, a hidden secret, a subset-sum encryption map, and a threshold decision over the decrypted residue. That is why the architecture split matters: one-hot MLPs and DeepSets-style nonlinear set aggregation solve the benchmark, while the tested small attention family does not.

The winning DeepSets-style circuit is related to the permutation-invariant architecture of Zaheer et al. [9]. We use that idea in a deliberately arithmetic way: lift local residue features nonlinearly, aggregate by summation, and classify the resulting pooled state. This result is narrower than a general claim about DeepSets and modular reasoning: this fixed-key decryptor has a structure that the local-lift plus sum pattern matches well.

Finally, the leakage analysis uses familiar interpretability tools. Linear probes follow the basic idea of training readouts on frozen internal features [10]. The local edit and patch-style controls are in the same family as activation patching and causal tracing methods [11], although this paper’s main patch result is a compact control rather than a full causal-circuit study. The follow-on late-code paper takes that causal question more directly.

8. Discussion

8.1 What The Paper Shows

The phrase “transformers can hide keys in weights” is too vague for this benchmark and largely misses the point.

The evidence is more specific:

- a neural model can learn a fixed-key modular decryptor in a clean non-trivial regime
- the right inductive bias matters a lot

- the winning computation looks like nonlinear local lifts plus additive set aggregation
- the learned solution is strongly tied to the training key
- leakage is asymmetric across attack surfaces
- that asymmetry includes a specific black-box/white-box split:
 - manifest-backed boundary attacks collapse the key to a tied rank-1 ambiguity class
 - literal cross-model secret and transcript recovery still fail
 - low-budget capability cloning still fails

That asymmetry is the core result. Black-box stealing is delayed. Boundary attacks collapse the key to an ambiguity class rather than a literal secret. Raw weights do not suddenly reveal the secret. But the penultimate state of a winning model already exposes a decryptor-equivalent representation almost linearly.

8.2 Scope

The scope is controlled on purpose. The benchmark is tiny, fixed-key, and deliberately clean; that is what makes the architecture split, rank-collapse result, and late-state leakage interpretable. Within that scope, the evidence supports a keyed-circuit story rather than a secure-vault story, a small-attention story, or a general re-keyability story.

8.3 Why The Leakage Story Matters

The most interesting outcome is that different notions of “recovery” separate:

- copying the function is possible, but not at tiny query budgets
- collapsing the secret to a narrow ambiguity class is easier than exact literal recovery
- probing the internal computation of one trained winner exposes decryptor state more easily than literal secret structure across many winners

That makes the benchmark useful as a controlled leakage microscope. The benchmark shifts attention from perfect hidden-key boxes to the form of leakage that appears first once a keyed circuit is learned.

9. Limitations

The strongest positive results live in tiny clean regimes, where the architecture split and leakage hierarchy are easiest to isolate. The main remaining limits are:

- The attention comparison is broad enough to be meaningful, but not exhaustive.
- Our chosen-query attacks show ambiguity-class collapse, but they are still not full cryptanalysis.

- The weight and activation meta-extraction results are negative only for the probes we tried.
- The positive same-secret fresh-A transfer result does not by itself turn this into a many-instance generalization paper.
- The local key-edit result and the negative cross-key retargeting result are both specific to the tested low-rank intervention path.
- A side shared-trunk plus tiny keyed-adaptor factorization probe also stayed weak; we keep it in Appendix A.6 rather than in the main text.

Two follow-on questions now look most natural. One is causal: decode the penultimate decryptor subspace, ablate it, and measure how the classifier collapses. The other is comparative: turn the leakage ordering in this benchmark into a dedicated follow-on paper about what leaks first and what remains stubbornly hard.

10. Conclusion

We asked whether neural circuits can learn fixed-key modular decryption, which architectures succeed, and what becomes extractable once they do.

Our answer is:

- yes, fixed-key modular decryption is learnable
- no, the tested small attention family is not the natural winner
- yes, the learned solution can be strongly key-bound
- and no, leakage does not reduce to one simple story

In our benchmark, the winning family is a nonlinear additive set circuit. Capability theft is possible but delayed. Boundary attacks can collapse the secret to a narrow ambiguity class without uniquely recovering it. The penultimate state of a trained winner almost linearly exposes the decryptor itself. Raw weights and fixed-query cross-model signatures, by contrast, are far less cooperative.

The result is more interesting than the original transformer-centric story. It suggests that the real science lives in the gap between learning a keyed function, emulating that function, and extracting the literal secret.

References

- [1] Oded Regev. “On lattices, learning with errors, random linear codes, and cryptography.” STOC 2005.
- [2] Zvika Brakerski, Adeline Langlois, Chris Peikert, Oded Regev, and Damien Stehle. “Classical Hardness of Learning with Errors.” STOC 2013.
- [3] Emily Wenger, Mingjie Chen, Francois Charton, and Kristin Lauter. “SALSA: Attacking Lattice Cryptography with Transformers.” IACR Cryptology ePrint Archive, 2022/935.

- [4] Cathy Li, Jana Sotakova, Emily Wenger, Mohamed Malhou, Evrard Garcelon, Francois Charton, and Kristin Lauter. “SALSA PICANTE: a machine learning attack on LWE with binary secrets.” ACM CCS 2023.
- [5] Cathy Yuanchen Li, Emily Wenger, Zeyuan Allen-Zhu, Francois Charton, and Kristin Lauter. “SALSA VERDE: a machine learning attack on Learning with Errors with sparse small secrets.” NeurIPS 2023.
- [6] Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. “Grokking: Generalization Beyond Overfitting on Small Algorithmic Datasets.” arXiv:2201.02177, 2022.
- [7] Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. “Progress measures for grokking via mechanistic interpretability.” ICLR 2023.
- [8] Andrey Gromov. “Grokking modular arithmetic.” arXiv:2301.02679, 2023.
- [9] Manzil Zaheer, Satwik Kottur, Siamak Ravanbakhsh, Barnabas Poczos, Ruslan Salakhutdinov, and Alexander Smola. “Deep Sets.” NeurIPS 2017.
- [10] Guillaume Alain and Yoshua Bengio. “Understanding intermediate layers using linear classifier probes.” arXiv:1610.01644, 2016.
- [11] Fred Zhang and Neel Nanda. “Towards Best Practices of Activation Patching in Language Models: Metrics and Methods.” arXiv:2309.16042, 2023.

Appendix A. Main Numerical Results

A.1 Clean $n=4$, $q=17$ Core Models

Model	Params	Best val	Test
Additive modulo	86	50.45%	49.86%
DeepSets modulo	64,257	99.95%	99.94%
One-hot MLP	27,649	99.87%	99.88%
Modulo transf. d1	216,865	50.41%	49.85%
Modulo transf. d2	410,129	50.45%	49.47%
Modulo transf. d4	796,657	50.22%	49.97%
Hybrid transf.	445,185	50.51%	49.80%

A.2 Key-Specificity Matrix

Src / Tgt	key 101	key 102	key 103
key 101	99.93%	50.03%	50.07%
key 102	49.93%	99.94%	50.68%
key 103	49.80%	50.36%	99.99%

A.3 Leakage Summary

Attack	Budget	Best outcome
Random stealing	200	50.23% test, 50.27% agr.
Random stealing	2,000	49.68%, 49.74%
Random stealing	20,000	50.44%, 50.49%
Random stealing	200,000	99.91%, 99.92%
Basis-window queries	68	0 / 3 exact recoveries
Binary-mask queries	255	0 / 3 exact recoveries
Boundary-manifest sweep	68 + local candidates	4 / 4 planted tied for rank 1, 0 / 4 exact
Weight probe zoo	28 ckpts	12.5% coord., 0% exact
Activation probe zoo	68 shared	0-12.5% coord., 0% exact
d-transcript probe zoo	68 shared	0.0-1.47% transcript pos., 0% exact
Structured stealing	68-255	49.58-50.36% test, 49.67-50.37% agr.

A.4 q / Width Boundary Sweep

q	Width	Test	Best val
17	128	99.61%	99.62%
17	256	99.85%	99.90%
23	128	66.94%	67.19%
23	256	99.85%	99.76%
31	128	50.33%	50.22%
31	256	74.43%	74.13%

A.5 Noise, Generalization, And Editability Controls

Control	Setting	Main outcome
Sigma ladder	q=17, width 256, sigma=0.10/0.25/0.50	99.85-99.88% test
Same-secret fresh-A transfer	held-out A_1, A_2	99.87%, 99.86%
Local key edit patch	rank 16, one-coordinate secret edit	93.83% edited-label test, 51.18% original-label test

Control	Setting	Main outcome
Activated patch retargeting	rank 16, 100 to 100,000 target examples	stays at chance

A.6 Shared Trunk + Tiny Keyed Adapter

We also tested a simple architectural split motivated by the idea that a shared reusable trunk might carry generic structure while a tiny adapter might carry the secret.

The result was weak.

A shared DeepSets trunk with adapter dimension 16, trained jointly across three keys, reached:

- 66.59% overall test accuracy with the learned adapters
- 55.30% overall test accuracy when the adapter was zeroed out

Per-key results were highly uneven:

- key 101: 99.67%
- key 102: 49.97%
- key 103: 50.15%

So this naive “generic trunk plus tiny secret patch” factorization did not cleanly solve the task. The secret-bearing computation does not obviously collapse into a reusable trunk plus a cheap keyed add-on, at least in this first version.